

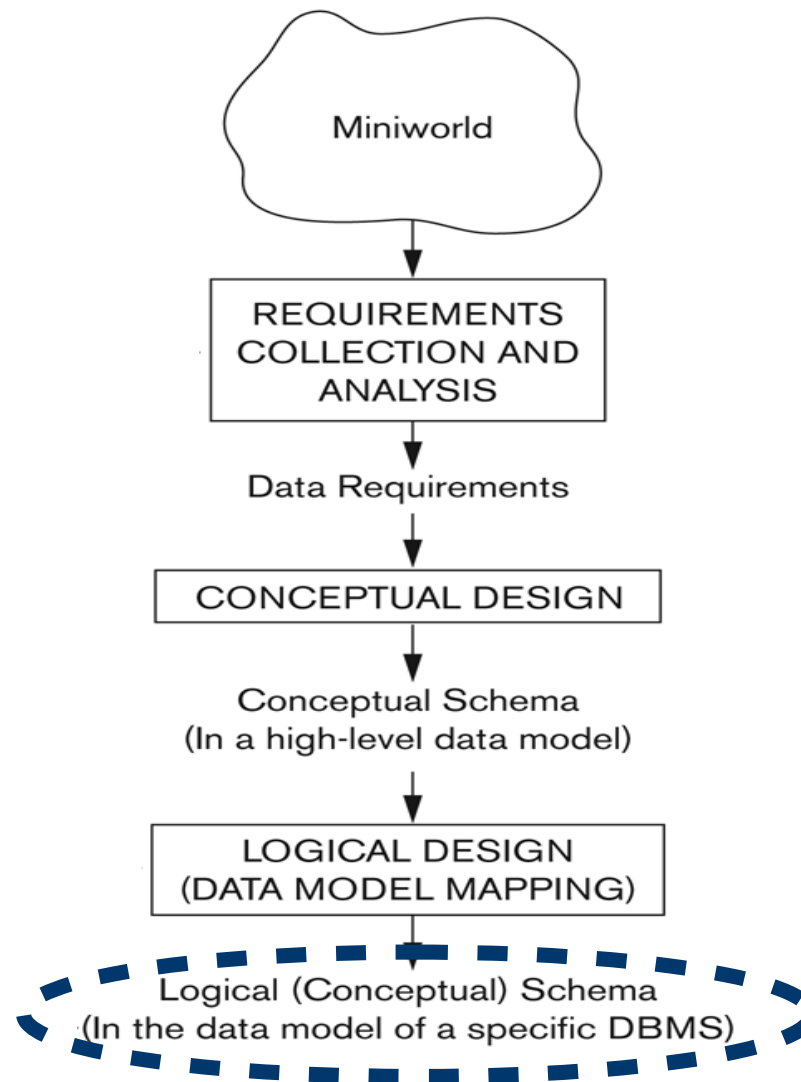
Database Technology

Topic 2: Relational Databases

Olaf Hartig

olaf.hartig@liu.se

Recall: DB Design Process



Relational Data Model

Relational Model Concepts

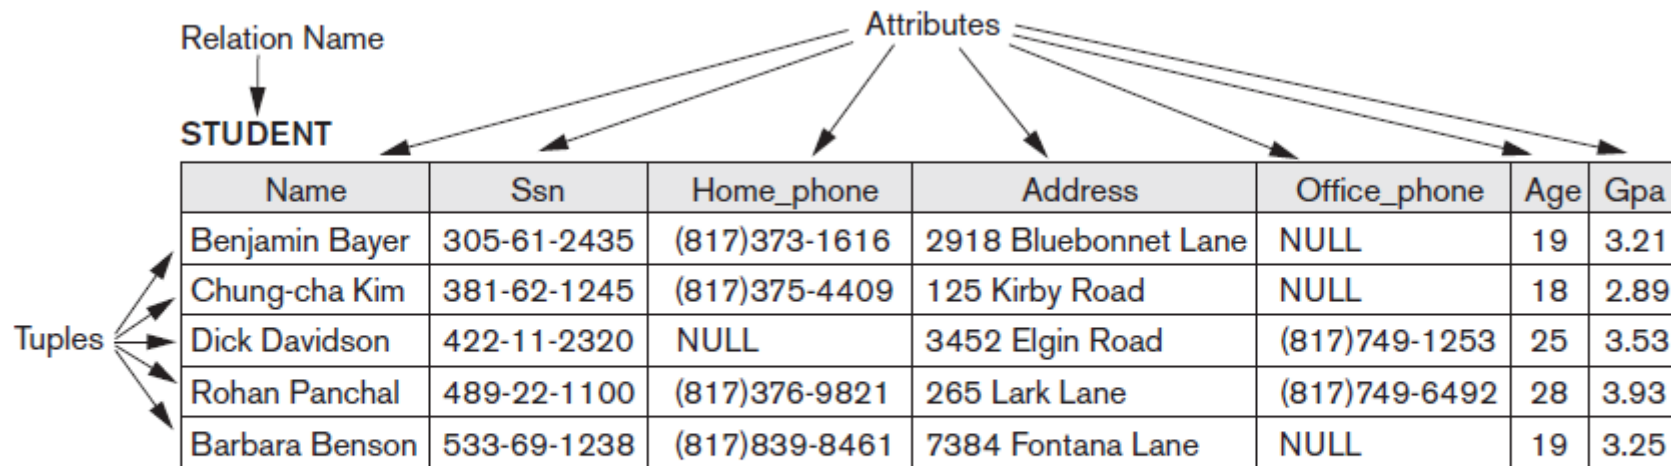
- Relational database: represent data as a collection of *relations*
 - Think of a relation as a table of values

Name	Ssn	Home_phone	Address	Office_phone	Age	Gpa
Benjamin Bayer	305-61-2435	(817)373-1616	2918 Bluebonnet Lane	NULL	19	3.21
Chung-cha Kim	381-62-1245	(817)375-4409	125 Kirby Road	NULL	18	2.89
Dick Davidson	422-11-2320	NULL	3452 Elgin Road	(817)749-1253	25	3.53
Rohan Panchal	489-22-1100	(817)376-9821	265 Lark Lane	(817)749-6492	28	3.93
Barbara Benson	533-69-1238	(817)839-8461	7384 Fontana Lane	NULL	19	3.25

- Each row (*tuple*) represents a record of related data values
 - Facts that typically correspond to a real-world entity or relationship
- Each column (*attribute*) holds a corresponding value for each row
 - Columns associated with a **data type** (**domain**)
 - Each column header: *attribute name*

Relational Model Concepts (cont'd)

- Relational database: represent data as a collection of *relations*
 - Think of a relation as a table of values



- **Schema** describes the relation
 - Relation name, attribute names and domains
 - Integrity constraints
- **Instance** (also called **state**) denotes the *current* contents of the relation
 - *Set* of tuples

Domains

- **Domain** is a set of *atomic* values
 - { 0, 1, 2, ... }
 - { Jo Smith, Dana Jones, Ashley Wong, Y. K. Lee, ... }
- **Atomic**: Each value indivisible
- Domains specified by **data type** rather than by enumeration
 - Integer, string, date, real, etc.
 - Can be specified by format
 - e.g., *(ddd)ddd-dddd* for phone numbers
(where *d* represents a digit)

Schemas and Attributes

- **Relation schema**

- A relation name R and a list of attributes $A1, A2, \dots, An$
- Denoted by $R(A1, A2, \dots, An)$

- **Attribute A_i**

- Name of a role in the relation schema R
- Associated with a domain **$\text{dom}(A_i)$**
- Attribute names do not repeat within a relation schema, but domains can repeat

- **Degree (or arity) of a relation**

- Number of attributes n in its relation schema

NULL Values

- Each domain may be augmented with a special value called NULL
 - Represent the values of attributes that may be unknown or may not apply to a tuple
 - If an attribute of a tuple is NULL, we cannot make any assumption about the value for that attribute (for that tuple)
- Interpretations for NULL values
 - Nothing is known about the value
 - Value exists but is (currently) not available
 - Value undefined (i.e., attribute does not apply to this tuple)
- For instance, Ashley's telephone number is NULL could mean
 - Ashley doesn't have a phone
 - Ashley has a phone but we don't know the number (perhaps withheld)
 - Ashley has a phone that has no number

Integrity Constraints

What are Integrity Constraints?

- Restrictions on the permitted values in a database instance / state
 - Derived from the rules in the miniworld that the DB represents
- 1. **Inherent model-based constraints** (also called **implicit constraints**)
 - Inherent in the data model, enforced by DBMS
 - e.g., duplicate tuples are not allowed in a relation
- 2. **Schema-based constraints** (also called **explicit constraints**)
 - Can be expressed in schemas of the data model, enforced by DBMS
 - e.g., films have only one director
 - Our focus here
- 3. **Application-based** (also **semantic constraints** or **business rules**)
 - Not directly expressed in schemas
 - Expressed and enforced by application program
 - e.g., this year's salary increase can be no more than last year's

Uniqueness Constraints

- Let R be a relation and K be a (sub)set of attributes of R
- If we specify the uniqueness constraint for K , then for any pair of tuples in R , the tuples must have a different value for at least one of the attributes in K
- Uniqueness must hold in all valid instances of R
- Uniqueness serves as a constraint on updates

Student

PN	FName	LName
19970218-1782	Jennifer	Li
19951223-6512	Paul	Smith
19990721-1222	Kim	Jonsson

Grade

Course	StPN	Grade
TDDD17	19970218-1782	4
TDDD43	19970218-1782	5
TDDD43	19951223-6512	3

Superkeys and Candidate Keys

- A set K of attributes of R is called a *superkey* of R if it has the
Uniqueness property: no two distinct tuples have the same values across all attributes in K
(i.e., we may define a uniqueness constraint for K)
- K is called a *key* of R if, additionally, it **also** has the
Minimality property: no proper subset of K has the uniqueness property
- Hence, every key is a superkey, but not every superkey is a key
- “candidate key” = key
 - used, in particular, if multiple different keys are possible

Primary Key

- There may be *more than one* candidate key in a relation
- **Primary key**: a particular candidate key is *chosen* as the primary
 - Diagrammatically, underline its attribute(s)
 - Tuples cannot have NULL for any primary key attribute
- Other candidate keys are designated as **unique**
 - Non-NULL values cannot repeat, but values may be NULL

Person1

PN	<u>Name</u>
19970218-1782	Jennifer
19970218-1782	Paul
19990721-1222	Jennifer

Person2

<u>PN</u>	Name
19970218-1782	Jennifer
19970218-1782	Paul
19990721-1222	Jennifer

Other Schema-Based Integrity Constraints

- **Entity integrity constraint:** No primary key value can be NULL
- **Domain constraint:** declared by specifying the datatype (domain) of the attributes
- **Referential integrity constraint**
 - see next slides

Referential Integrity Constraints (Motivation)

- Consider the following two relations

Student

<u>PN</u>	Name
19970218-1782	Jennifer
19951223-6512	Paul
19990721-1222	Kim

Grade

<u>Course</u>	<u>StPN</u>	Grade
TDDD17	19970218-1782	4
TDDD43	19970218-1782	5
TDDD43	19951223-6512	3

- We may want to make sure that for every student for which we record grades (in the Grade relation) we have a record in the Student relation
- That is, assuming the given instance of the Student relation, it would be invalid to have the following tuple in the Grade relation:

(TDDD17, 20010219-6678, 4)

Referential Integrity Constraints

- Maintains consistency among tuples in two relations
- Allows every tuple in one relation to refer to a tuple in another
- Formally:
 - Let PK be the primary key in a relation $R1$
 - e.g., $PK = \{ PN \}$ in the Student relation on the previous slide
 - Let FK be a set of attributes for another relation $R2$
 - e.g., $FK = \{ StPN \}$ in the Grade relation on the previous slide
 - The attribute(s) FK have the same domain(s) as the attribute(s) PK
 - Constraint: For every tuple $t2$ in $R2$, either
 - i) there is a tuple $t1$ in $R1$ such that the value that $t1$ has for PK is the same as the value that $t2$ has for FK , or
 - ii) the value that $t2$ has for FK is NULL
 - e.g., for every tuple $t2$ in the Grade relation, there is a tuple $t1$ in the Student relation such that the PN value of $t1$ is the same as the StPN value of $t2$, or the StPN value of $t2$ is NULL

www.liu.se